


# Uber

## Business Case Study On Application of Big Data Technologies

► Business Case Study Presentation  
by:

- Divya Vemula
- Rahul Chauhan
- Mani Krishna Tippani
- Shaheryar Nadeem




# Introduction to Uber and Big Data

## ► Overview of Uber

- Founded in 2009, Uber operates in ride-sharing, food delivery (Uber Eats), and freight logistics.
- Uber operates in more than 60 countries and has millions of riders and drivers worldwide, making it one of the largest transportation networks globally.
- Uber generates revenue from ride-sharing, food delivery, freight logistics, and advertising. It operates using a commission-based model, taking a percentage from both drivers and businesses.

## ► Why Big Data for Uber?

- Uber generates 15 TB of data daily, processes it at a real-time rate of millions of events per second, and stores over 100 petabytes of data.
- Massive scale operations demand efficient data handling.
- Real-time decision-making crucial for rider-driver matching, dynamic pricing, and operational optimization.



# Challenges at Uber and Need for Hadoop

## Data Infrastructure challenges:

- Scalability became a challenge as Uber's operations expanded globally. As Uber handles petabytes of data daily, total incoming data volume is growing at an exponential rate Replication factor & several geo regions copy data.
- the infrastructure must scale seamlessly without re-architecture, while keeping costs low due to Uber's low-margin business model
- Data must be processed within seconds of generation to support real-time demand-supply adjustments and dynamic features like surge pricing.
- Uber relied on relational databases like MySQL and PostgreSQL. Uber needs to offer both **SQL** and **programmatic interfaces** to accommodate users with varying technical skills and support diverse business needs.
- Data silos created difficulties in generating a unified, global operational view.

## Operational Bottlenecks:

- Data processing latency exceeded 24 hours, hindering timely decision-making.
- The inability to handle rapidly growing data volumes constrained operational efficiency and innovation.

## Uber's Data Challenges Before 2014:

- Before 2014, Uber used a few traditional OLTP databases (MySQL, PostgreSQL) to manage data., with no global view of the data
- As Uber expanded, data volume and complexity increased, requiring a centralized data warehouse to support analysis.

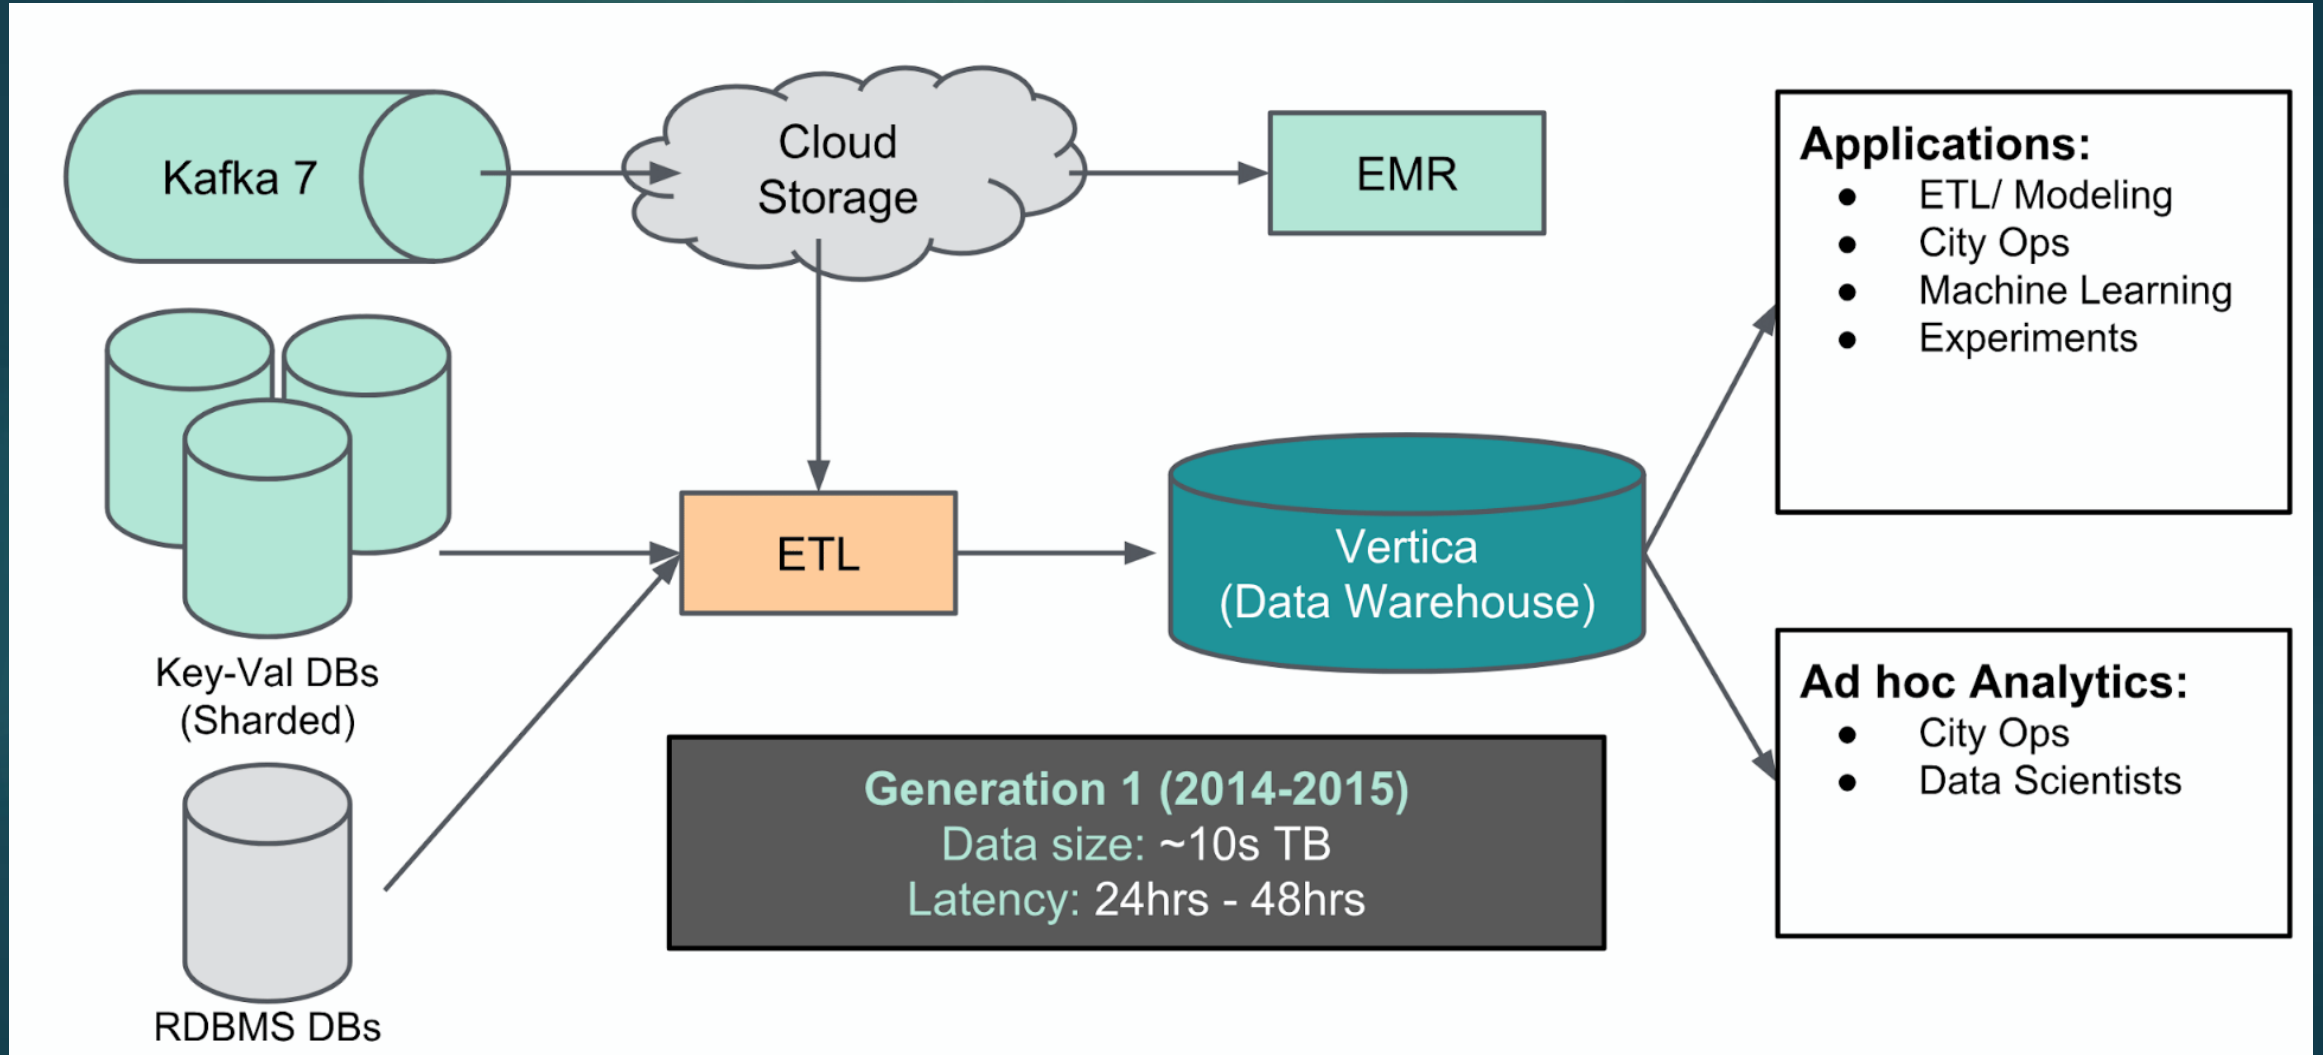
## Beginning of Big Data at Uber 2014-15

- The Beginning of Uber's Big Data platform allowed to aggregate all of Uber's data in one place i.e., Vertica DW and provide standard SQL interface for users to access data.

## Limitations

- ETL jobs lacked schema enforcement and formal contracts between data producers and consumers, which led to duplicate data ingestion.
- Most of data was in JSON format, and ingestion jobs were not resilient to changes in the producer code.

- Ad hoc ETL processes, caused data reliability concerns, inefficiencies and difficulty in integrate new data types.
- This led to pressure on upstream data sources and increased storage, operational costs.

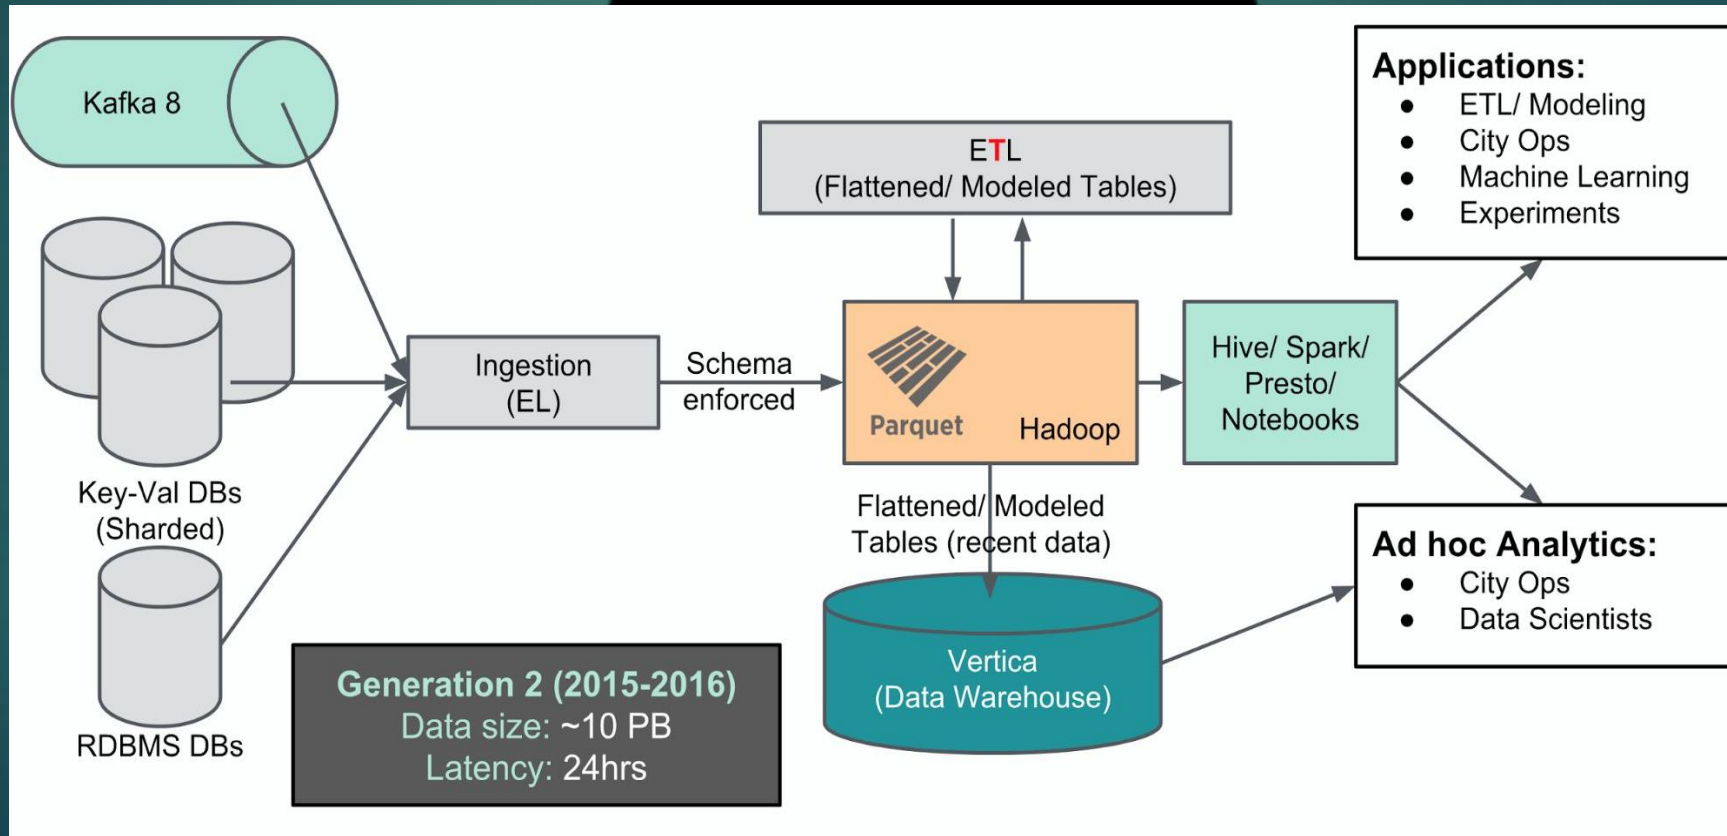


# Arrival of Hadoop in Uber (2015-16)

Uber transitioned into a Hadoop data lake for raw data ingestion, avoiding transformations during ingestion to reduce pressure on online datastores.

## Data access data in Hadoop:

- **Presto** - enable interactive ad hoc user queries,
  - **Apache Spark** - facilitate programmatic access to raw data
  - **Apache Hive** - workhorse for extremely large queries.
- Data modeling and transformations only in **Hadoop**, enabling fast backfilling & recovery when issues arose.
  - Transition from JSON to **Apache Parquet**, to store both schema and data together, ensuring data consistency and preventing issues from upstream format changes.

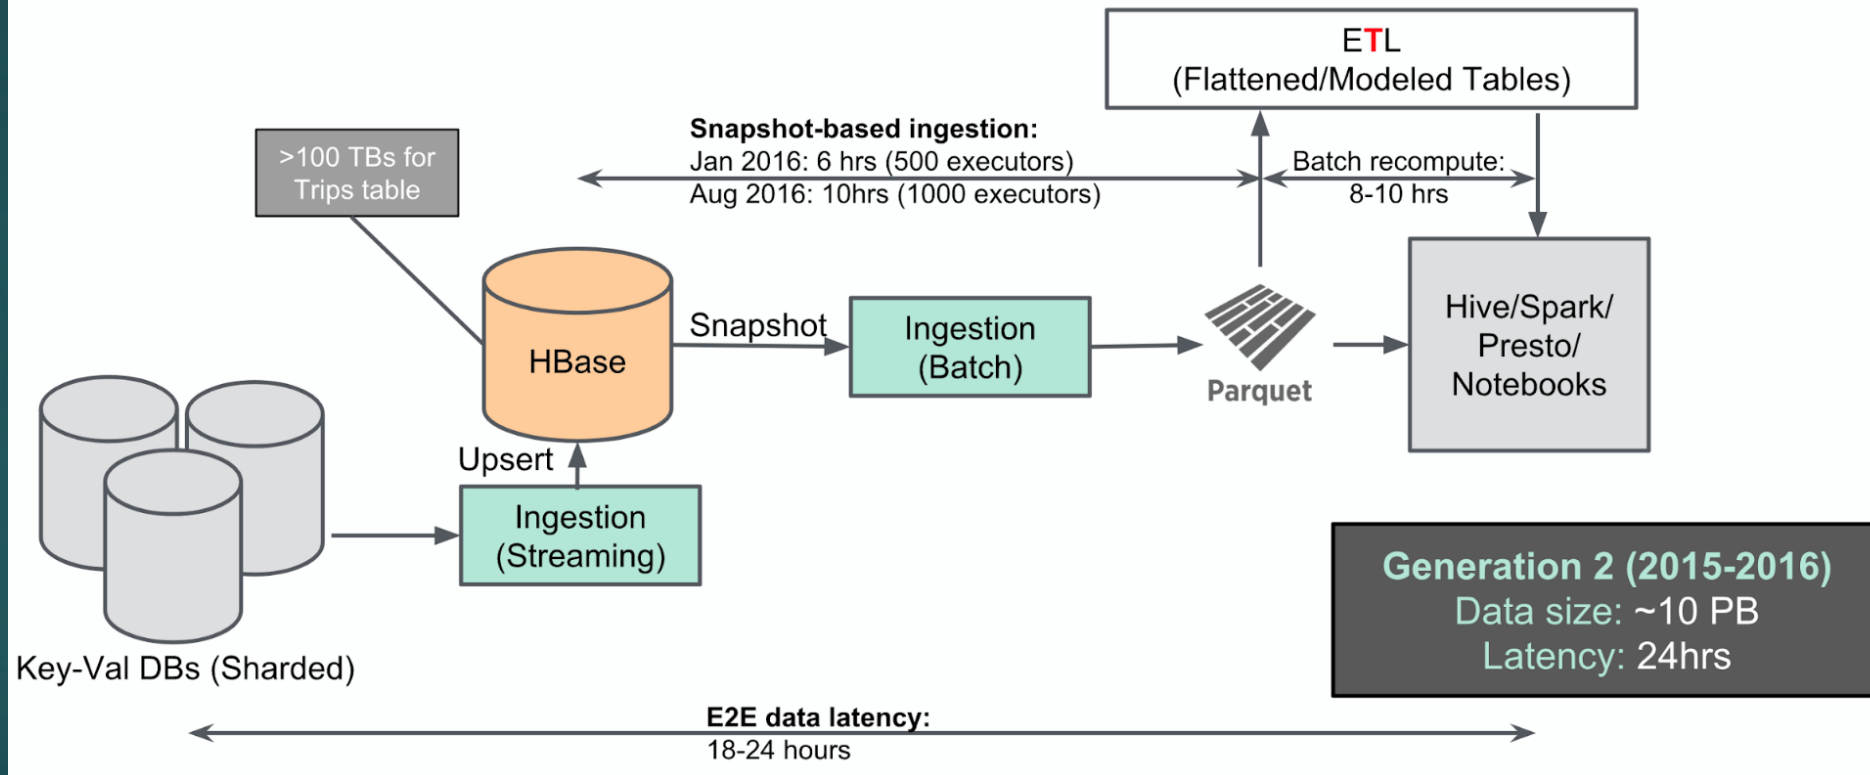




# Limitations

- **Data latency** remained a challenge, with new data becoming available only **every 24 hours**, hindering real-time decisions.
- ETL and modeling processes were still bottlenecks due to the need to **recreate entire datasets** during each run.
- **Snapshot-based ingestion** was required for data updates, which led to **over 1,000 Spark executors** running jobs that could take **20+ hours**.

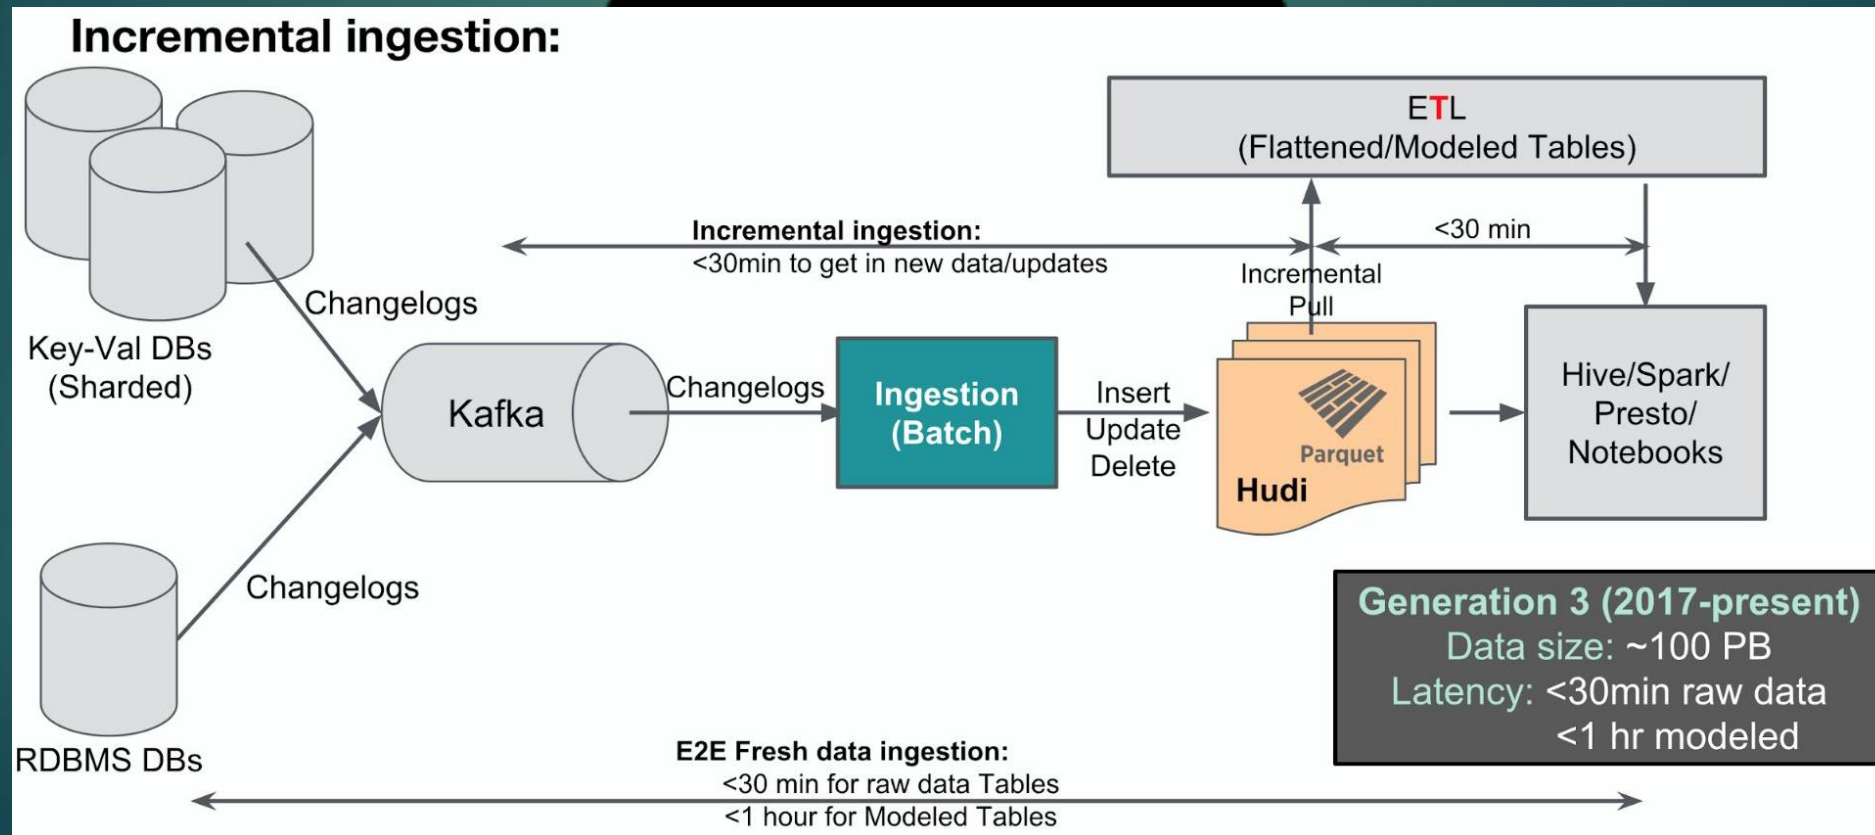
## Why does data latency remain at 24 hours?



The latency for new data was still over one day, a lag due to the snapshot-based ingestion of large, upstream source tables that take several hours to process.

# Rebuilding Big Data platform of Uber (2017- Present)

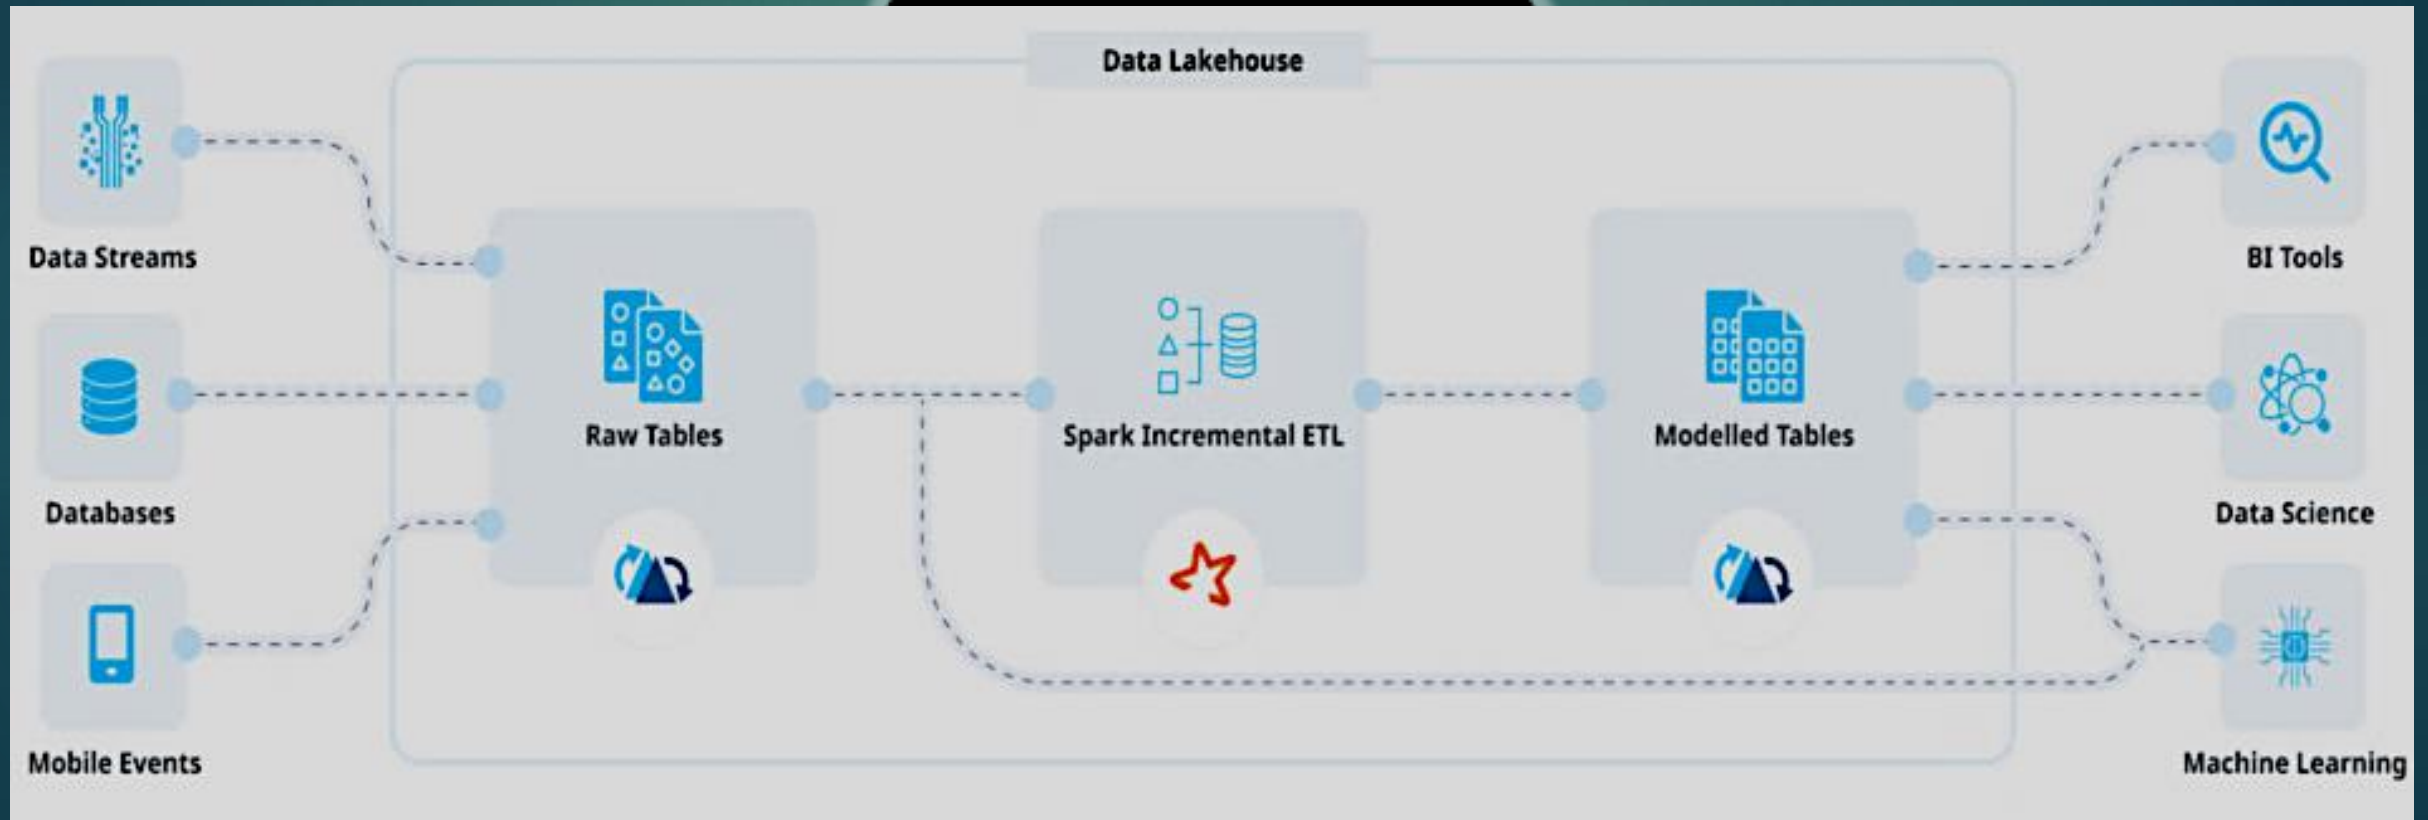
- **Hadoop and Parquet** lacked support for **updates and deletes**, requiring a solution for efficient data modifications.
- To address above issue, **Hudi** (Hadoop Upserts and Incremental) is **developed by Uber**, an open-source **Spark** library for **incremental ingestion**, updates, and deletes is used.
- With **Hudi**, data latency reduced from 24 hours to **less than one hour**, improving **ETL jobs** and enabling real-time analytics.
- The **Marmaray** platform ingests these events in mini-batches, applying changes to Hadoop using **Hudi** and also avoids inefficient transformations during raw data ingestion, using an **EL** (Extract-Load) model instead of traditional ETL.





- Uber uses incremental data processing to ensure data freshness and cost-efficiency at scale.
- **Apache Hudi**, supports incremental updates by handling late-arriving data seamlessly, Combines stream processing real-time data with the efficiency of batch processing.
- Incremental processing reduces the computational costs and duration of batch jobs, allowing for more frequent updates.

For example, Late-arriving updates, such as rider tips added after trips, are processed incrementally, avoiding recalculation of entire partitions.



- Incremental processing introduces mutable, database-like features in data lakes, enhancing performance and reliability.
- Data from various sources (e.g., Kafka, databases) is continuously ingested into raw datasets.
- Ingestion uses upserts to add or update data in raw datasets incrementally without overwriting entire partitions.

### Incremental Pull and Joins:

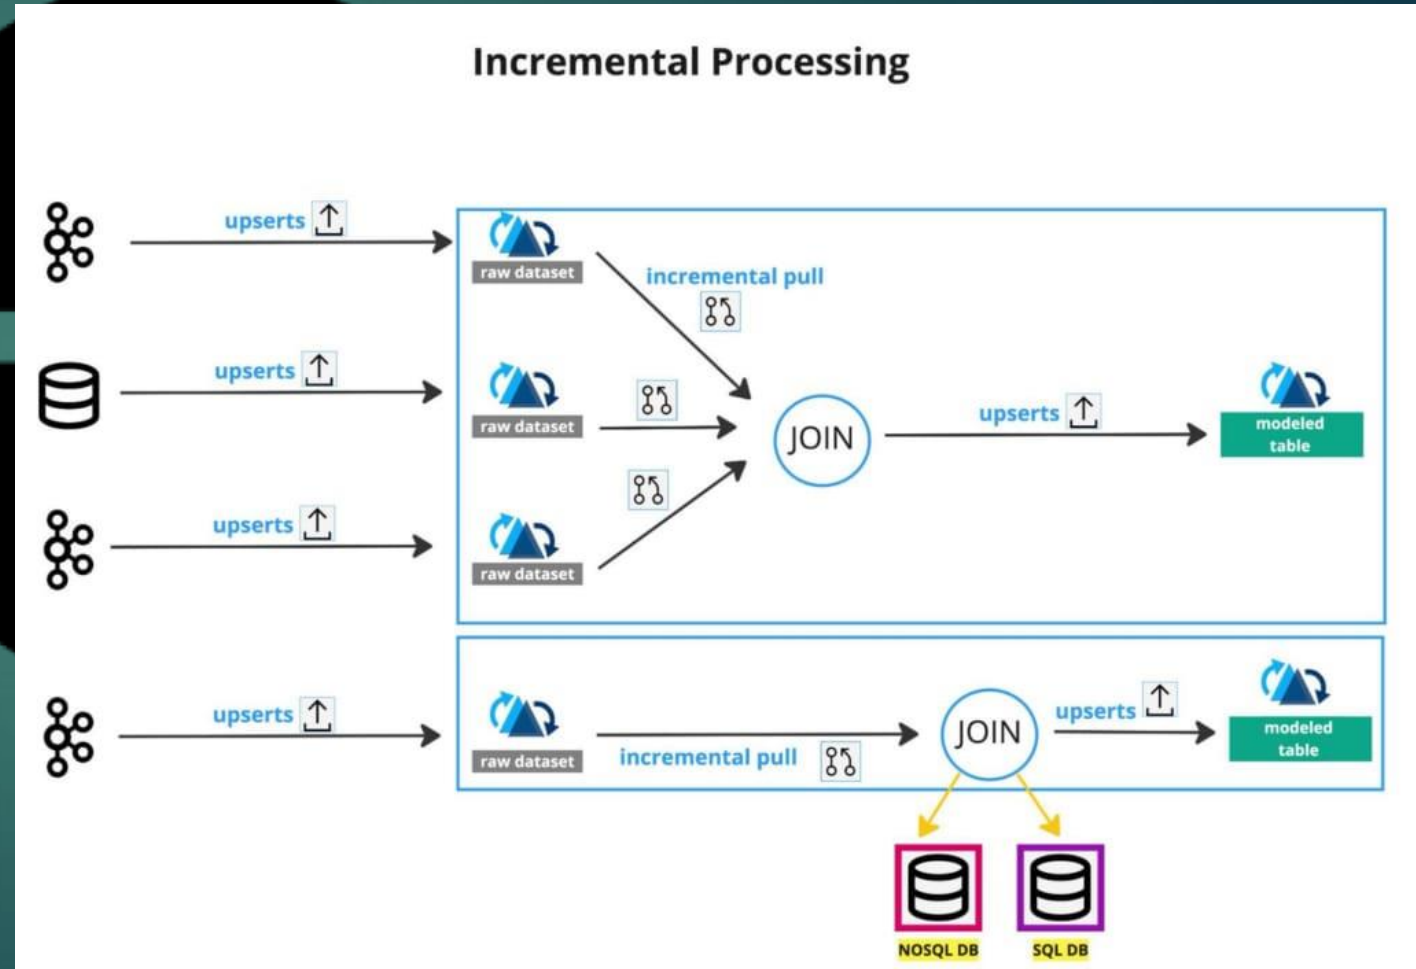
- Incremental pull operations fetch only the newly ingested or updated data from raw datasets.
- These incremental updates are combined using joins to integrate data from multiple raw datasets.

### Modeled Table:

- Results are written to modeled tables using upserts.
- Modeled tables have processed and aggregated datasets for downstream applications.

### Database Integration:

- Incremental data is also joined with external databases to enhance the modeled table with additional information in required workflows.



## Performance and Cost Savings

Transitioning to Apache Hudi enabled strongly consistent replication by moving only incrementally changed files using Hudi metadata, avoiding inconsistencies with plain parquet tables.

- Incremental ETL pipelines reduced **pipeline run time by up to 82%**, achieving a **50% average reduction** compared to batch ETL.
- Significant cost reductions, with **up to 78.6% savings** on processing costs for dimensional driver tables.
- Memory and CPU usage improved by **73% and 59%, respectively** for key pipelines.

| Pipeline                                    | vcore_seconds | memory_seconds | Cost    | Run Time (mins) |
|---|---------------|----------------|---------|-----------------|
| Batch ETL of Dimensional Driver Table       | 3,129,130     | 23,815,200     | \$11.39 | 220             |
| Incremental ETL of Dimensional Driver Table | 1,280,928     | 6,427,500      | \$2.44  | 39              |
| Difference                                  | 1,848,202     | 17,387,700     | \$8.95  | 181             |
| % Improvement                               | 59.06%        | 73.01%         | 78.57%  | 82.27%          |
| Batch ETL of Driver Status Fact Table       | 2,162,362     | 5,658,785      | \$3.30  | 94              |
| Incremental ETL of Driver Status Fact Table | 1,640,438     | 3,862,490      | \$2.45  | 48              |
| Difference                                  | 521,924       | 1,796          | \$0.85  | 46              |
| % Improvement                               | 24.13%        | 31.74%         | 25.75%  | 48.93%          |

# Big Data Prospects and Developments in Uber

- Uber is migrating its batch data analytics and ML training stack to Google Cloud Platform (GCP) to enhance productivity, cost efficiency, and data governance.
- The strategy involves initially replicating the on-prem stack on GCP's IaaS, with plans to adopt PaaS offerings like Dataproc and BigQuery for scalability and performance.
- The migration strategy leverages cloud's object storage with HDFS compatibility and translating HDFS paths to object store paths via a "Path Translation Service."
- Uber's existing container and deployment systems, built to work across on-prem and cloud, support the expansion of batch data ecosystem microservices to GCP IaaS without major modifications.
- Extending HiveSync to support active-active bi-directional replication and ongoing updates to the cloud-based data lake.

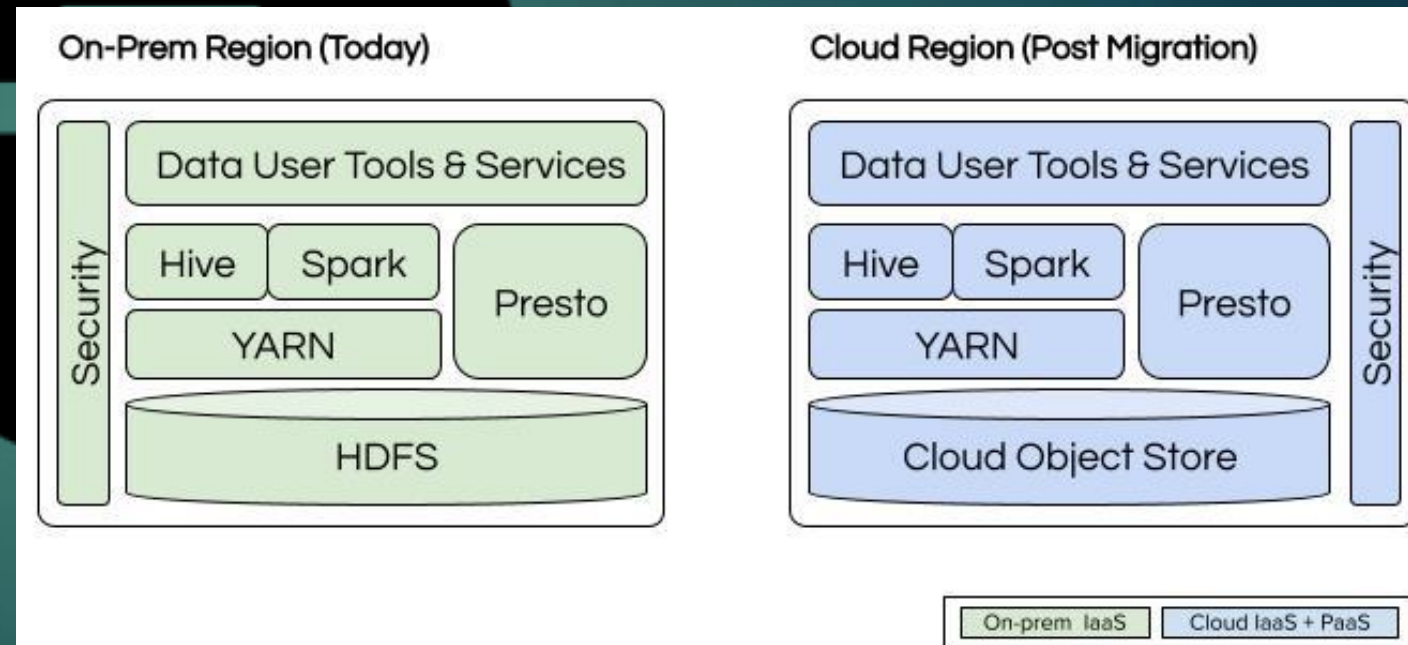


Figure 2: Uber Batch Data Stack: Pre and Post Migration



# References

- [Uber's Big Data Platform: 100+ Petabytes with Minute Latency](https://www.uber.com/blog/uber-big-data-platform/)

(<https://www.uber.com/blog/uber-big-data-platform/>)

- [Cost-Efficient Open Source Big Data Platform at Uber](https://www.uber.com/blog/cost-efficient-big-data-platform/)

(<https://www.uber.com/blog/cost-efficient-big-data-platform/>)

- [Setting Uber's Transactional Data Lake in Motion with Incremental ETL Using Apache Hudi](https://www.uber.com/blog/ubers-lakehouse-architecture/)

(<https://www.uber.com/blog/ubers-lakehouse-architecture/>)

- [Modernizing Uber's Batch Data Infrastructure with Google Cloud Platform](https://www.uber.com/blog/modernizing-ubers-data-infrastructure-with-gcp/#:~:text=Uber%20runs%20one%20of%20the%20largest%20Hadoop%20installations%20in%20the%20world.&text=We%20will%20leverage%20open%20standards%20such%20as,the%20teams%20within%20the%20data%20platform%20organization.)

([https://www.uber.com/blog/modernizing-ubers-data-infrastructure-with-](https://www.uber.com/blog/modernizing-ubers-data-infrastructure-with-gcp/#:~:text=Uber%20runs%20one%20of%20the%20largest%20Hadoop%20installations%20in%20the%20world.&text=We%20will%20leverage%20open%20standards%20such%20as,the%20teams%20within%20the%20data%20platform%20organization.)

[gcp/#:~:text=Uber%20runs%20one%20of%20the%20largest%20Hadoop%20installations%20in%20the%20world.&text=We%20will%20leverage%20open%20standards%20such%20as,the%20teams%20within%20the%20data%20platform%20organization.](https://www.uber.com/blog/modernizing-ubers-data-infrastructure-with-gcp/#:~:text=Uber%20runs%20one%20of%20the%20largest%20Hadoop%20installations%20in%20the%20world.&text=We%20will%20leverage%20open%20standards%20such%20as,the%20teams%20within%20the%20data%20platform%20organization.))

- [DataMesh: How Uber laid the foundations for the data lake cloud migration](https://www.uber.com/blog/datamesh/)

(<https://www.uber.com/blog/datamesh/>)

- [DataCentral: Uber's Big Data Observability and Chargeback Platform](https://www.uber.com/blog/datacentral-ubers-observability-and-chargeback-platform/?uclid=bb4c38b1-fe21-44a3-be95-f97578d04874)

(<https://www.uber.com/blog/datacentral-ubers-observability-and-chargeback-platform/?uclid=bb4c38b1-fe21-44a3-be95-f97578d04874>)